

22/11/21

Transaction Processing

A transaction is a unit of a program that possibly access and update different data items present in a database in a all or none manner.

Banking Database

ACID Properties of Transaction Process :-

- (1) A : Atomicity
- C : Consistency
- I : Isolation
- D : Durability

Atomicity: Either all operation of the transaction are reflected properly in database or none.

Consistency:

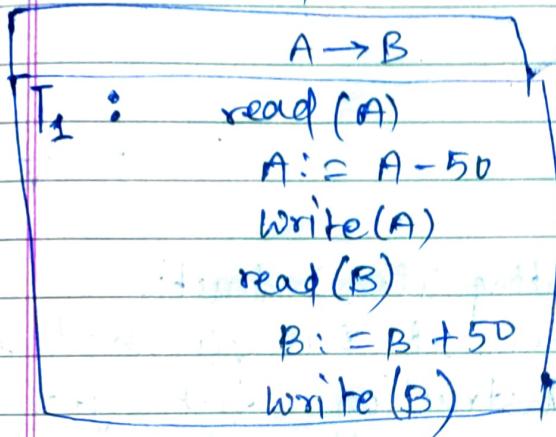
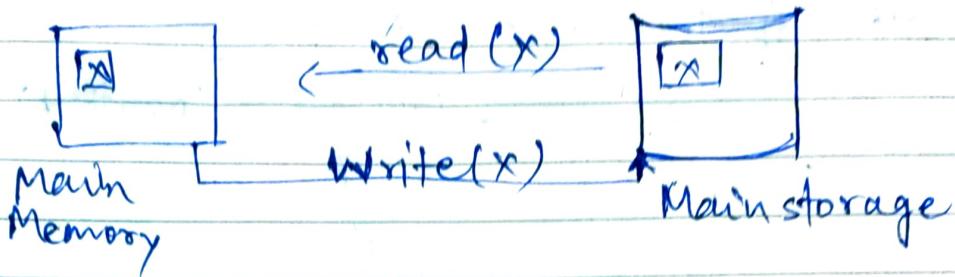
Isolation: When More than one transaction are processed concurrently.

Durability:

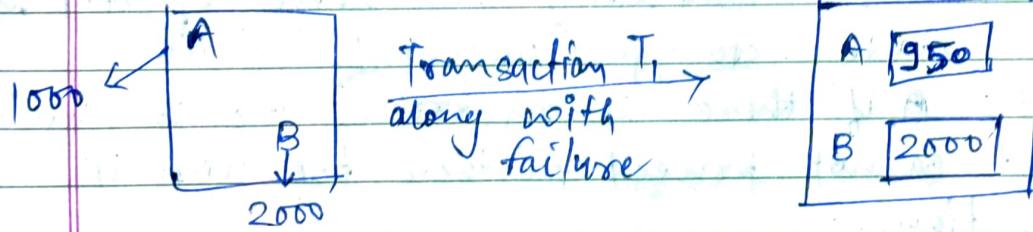
A Simple Transaction Model:

① read(x)

② write(x)



① Atomicity :-



② Consistency :- The sum of the balance should before and after transaction must be same.

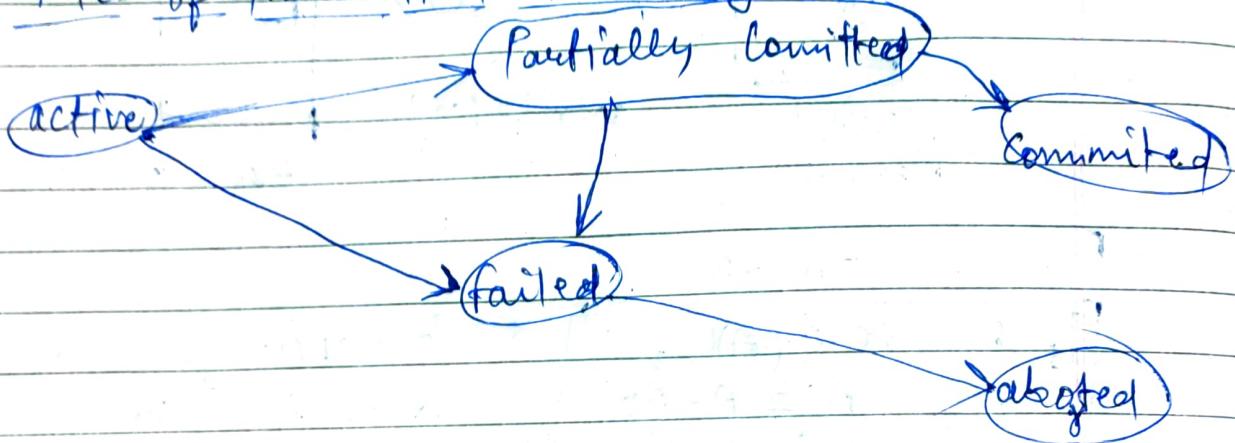
③ Durability

Volatile Storage

Non-Volatile Storage

Stable storage

States of Transaction Processing:-



active: when it is executing its instruction.

Partially Committed: when it has finished its last instruction but yet to update the data items in the stable storage.

failed:

Committed: When data items has been updated in stable storage (Successful Execution)

Failed: If there is an ~~error~~ due to which it cannot proceed further for normal execution.

8

(1) logical error

(2) Hardware/ Network failure

(3) Security issues

Aborted: When its effect has been rolled back.
(Unsuccessful Execution)

A transaction is said to be terminated if it has either committed or aborted.

Concurrency :

(1) Better throughput / Better resource utilization

Throughput : Number of transaction completed per unit time.

(2) Reduce waiting time :-

~~T2~~ T2: read(A)

temp := A * 0.1

A := A - temp

write(A)

read(B)

B := B + temp

write(B)

Serial Schedule :

T₁

read(A)

A := A - 50

write(A)

read(B)

B := B + 50

write(B)

Commit

T₂

Schedule - I

read(A)

temp := A * 0.1

A := A - temp

write(A)

read(B)

B := B + temp

write(B)

- Commit

Ques How to generate concurrent schedule?

Schedule - II

	T_1	T_2
$A \leftarrow 1000$		
$B \leftarrow 2000$		
	read(A) $\rightarrow 1000$	
	$A := A - 50$	
	Write(A) $\rightarrow 950$	
		read(A) $\rightarrow 950$
		$95 \leftarrow \text{temp} := A \times 0.1$
		$A := A - \text{temp}$
		Write(A) $\leftarrow 855$
	read(B) $\rightarrow 2000$	
	$B := B + 50$	
	Write(B) $\rightarrow 2050$	
	Commit	
		read(B) $\rightarrow 2050$
		$B := B + \text{temp}$
		Write(B) $\rightarrow 2145$
		Commit
	final value of $A = 855$	
	... or $B = 2145$	
	Total Val = 3000	

- Schedule II is a concurrent schedule which preserves the consistency property.

Ques Is it the case that every concurrent schedule maintains the consistency property?

Schedule: III

Date _____
Page No. _____

T₁

read(A) → 1000

A := A - 50

Write(A) → 950

read(B) → 2000

B := B + 50

Write(B) → 2050

Commit

T₂

read(A) ← 1000

temp = A × 0.1

A := A - temp

Write(A) → 900

read(B) → 2000

B := B + temp

Write(B) → 2100

Commit

(2) final value of A = 950
final value of B = 2100
sum = 3050

∴ Schedule III leads to the database to an inconsistent state.

Ques. Can we make a concurrent schedule which is equivalent to a serial schedule??

Serializable Schedule

How to determine whether two schedule are equivalent?

Schedule II

T_i
read(A)
write(A)

read(B)
write(B)
comes

T_j
read(A)
write(A)

read(B)
write(B)
comes

Conflict - serializability :-

T_i
read(A)

T_j
write(A)

Case:I T_i 's statement is scheduled first and the T_j 's statement is later

Case:II T_j 's statement is scheduled first and T_i 's statement is later

T_i T_j

T_Q

database

$T_i : \text{read}(A)$	$T_j : \text{read}(A) \rightarrow \text{No Conflict}$
$T_i : \text{read}(A)$	$T_j : \text{write}(A)$
$T_i : \text{write}(A)$	$T_j : \text{read}(A)$
$T_i : \text{write}(A)$	$T_j : \text{write}(A)$

↓
Conflict

We say that two instructions are non-conflicting if the order of execution does not make any difference.

This is only possible when both are read statements.

~~S~~ → → →

swapping means changing the order of execution of two non-conflicting statement.

A schedule is said to be conflict serializable if we can obtain a serial schedule by a series of swap operations between two non-conflicting statements.

- (1) If T_i and T_j operates on two different data items.
- (2) If T_i and T_j operates on same data items but both of them are read statement.

Schedule II

T_1

read(A)

write(A)

T_2

read(B)

write(B)

T_1

read(A)

write(A)

read(B)

write(B)

T_2



read(B)
write(B)

swapping two statement

T ₁	T ₂	T ₁	T ₂
read(A)		read(A)	
write(A)		write(A)	
←	read(A)	read(B)	
← read(B)	→	write(A)	
write(B)		write(B)	
	read(B)		read(B)
	write(B)		write(B)

From schedule II by a series of swap operation we are able to reach to a serial schedule. Hence schedule II is conflict serializable.

Schedule III

T ₁	T ₂	
read(A)		
	read(A)	
	write(A)	
←	read(B)	
← write(A)	→	
read(B)		
write(B)		
	write(B)	

Schedule III is not conflict serializable.

29/11/20

Date	
Page No.	

- If there are two transaction T_i and T_j working on two different data items P (by T_i) and Q (access by T_j) Then there absolutely no conflict.
- If $-- T_i$ and T_j --- one data item P and operation is read (R) then there no conflict.
- If $-- T_i$ and T_j --- one data item P and one of operation is write (W) then there is conflict.

If a schedule s can be transformed into a schedule s' by a series of swap operation of non-conflicting statements then we call s and s' are conflict equivalent.

~~if~~ A schedule s is said to be conflict serializable if it is equivalent to a serial schedule.

Ques. Given a schedule How to test whether it is conflict serializable or not?

Solution:

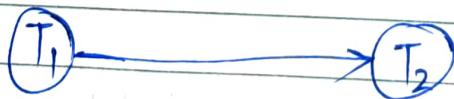
Consider for the given Schedule s we construct a directed graph $G_1(V, E)$:: Precedence Graph

$V(G) \rightarrow$ set of Transaction
 $V(G) = \{T_1, T_2, \dots, T_m\}$

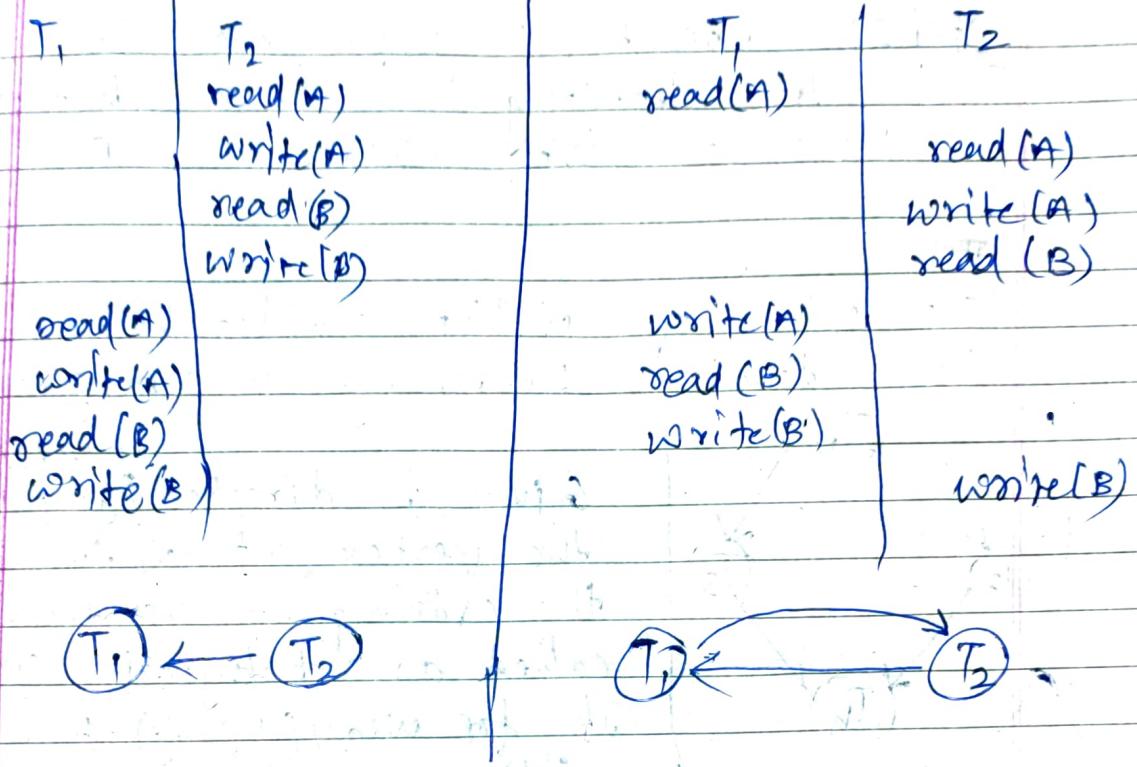
① In G_i there is a directed edge $T_i \rightarrow T_j$ ~~for~~. if any one of the following three cases arrives:

- (A) T_i executes a write(α) statement before T_j executes a read(α) statement.
- (B) T_i - - - a read(α) - - - before
 ~~T_j~~ T_j - - a write(α) - - - .
- (C) T_i - - - a write(α) - - - before
 ~~T_j~~ T_j - - - a write(α) - - - .

~~Ques~~ T_1 | T_2] for this schedule
read(A)
write(A)
read(B)
write(B)
read(A)
write(A)
read(B)
write(B)



~~Ans~~



Time requirement will be proportional to number of data items involved.

If the precedence graph contains a cycle then the schedule is not conflict serializable.

How to detect a cycle in a directed graph?

Graph Traversal

$$G(V, E)$$

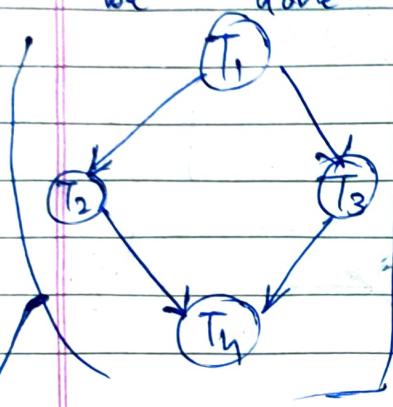
Test for conflict serializability :-

Step 1: Construct the precedence graph:

Step 2: Execute a graph traversal algorithm for detecting there is a closed cycle or not.

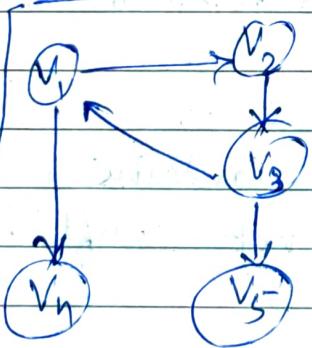
#

The order of conflict serializability can be done by performing by topological sorting.



Given a directed [↑]graph with the vertex set $\{v_1, v_2, \dots, v_n\}$, the ordering of vertices is called a topological sorting, if for every $(v_i, v_j) \in E(v)$

$$P(v_i) < P(v_j)$$



for directed cyclic graph it cannot have a topological sorting.

Topological Sorting : $T_1 \rightarrow T_2 \rightarrow T_3 \rightarrow T_4$

$T_1 \rightarrow T_3 \rightarrow T_2 \rightarrow T_4$

Such schedule are called non-recoverable schedules.

We will not let T_2 Commit unless T_1 commits.

T_1
read(A)
write(A)

T_2
read(A)
write(A)
Commit

dead(B)



T_1	T_2	T_3
read(A)		
read(B)		
write(A)		
	read(A)	
	write(A)	
		read(A)
(abort)	(abort)	(abort)

The phenomena in which a single transaction failure leads to a series of transaction roll backs. It is called cascading-roll back.

Cascadeless Schedule:

A " " is one where each pair of transaction T_i and T_j such that T_j reads a data item previously written by T_i , the commit operation of T_i appears before the read option of T_j .

\$
\$

\$
\$

- Multiple transaction are getting executed + transaction may fail.

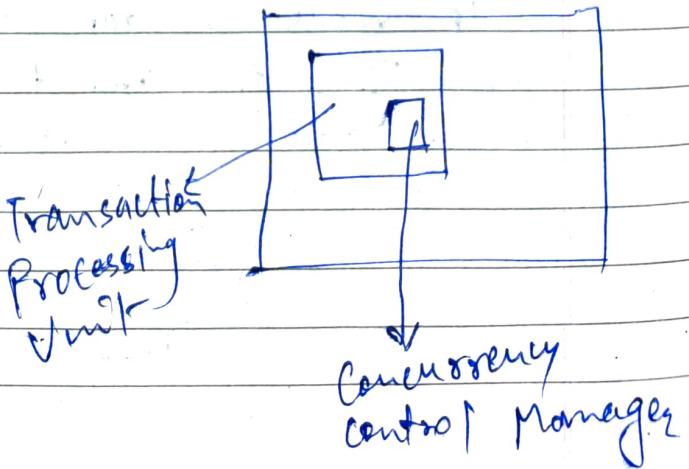
Q How to maintain concurrency in presence of transaction failure?

- Two phase locking Protocol
- Snapshot Isolation

Lock Based Protocols

① Shared mode lock : If a transaction T_i has obtained a shared mode lock (denoted by S) on data item & then ~~implication~~ T_i can only read the data item but cannot write on it.

② Exclusive Mode lock :- If a transaction T_i has obtained a exclusive mode lock (denoted by X) then T_i can both read and write.

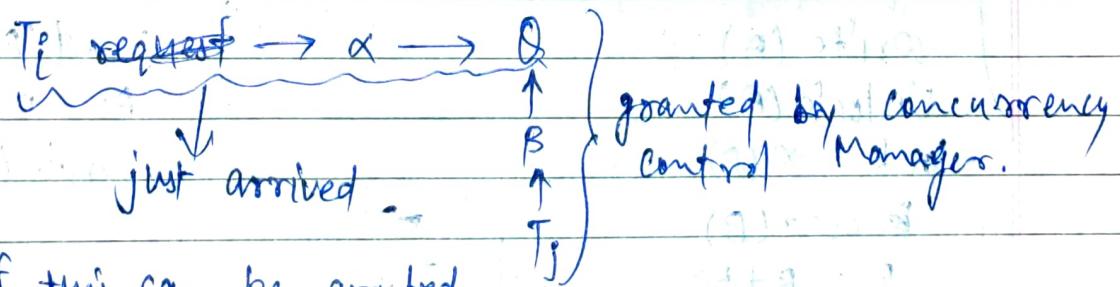


$T_i \rightarrow$ write operation
database (D)

- T_i has to make a request to the concurrency control Manager for having a lock in appropriate mode (exclusive mode).
- Upon having the request the concurrency control Manager either may grant it or may not.
why?

Given a set of lock modes we can define the notion of compatibility.

- ① Let α and β two arbitrary lock modes.



If this can be granted immediately even if

The lock modes α and β are compatible with each other.

T_i	T_j
S	S
S	True false
X	false false

The compatibility of locks can be represented in the form of a boolean matrix.

T_i

lock - S(Q) : T_i is requesting a shared mode lock on the data item Q.

lock - X(Q) : T_i is requesting a exclusive mode lock on the data item Q.

Unlock (Q) : A transaction T_i releases its lock on Q.

T₁ : Transferring 50 rs/- Acc A to B. ^{from}

T₁ { lock - X(A)

read (A)

A := A - 50

write (A)

Unlock (A)

lock - X (B)

B: read (B)

B := B + 50

write (B)

Unlock (B)

commit

} Read (A)

A := A - 50

write (A)

read (B)

B := B + 50

write (B)

T₂

lock - S(A)

read (A)

Unlock (A)

lock - S(B)

read (B)

Unlock (B)

display (A+B)

Commit

2/2) e)

Date _____
Page No. _____

Q. A possible schedule for concurrent execution of T_1 & T_2 :

T_1
lock - X(A)

read(A)
 $A := A + 50$

write(A)
Unlock(A)

lock - S(A)

read(A)
Unlock(A)

lock - S(B)

lock - X(B)

read(B)

$B := B + 50$

Write(B)

Unlock(B)

read(B)
Unlock(B)
display(A+B)

Concurrency Control / Manager

grant - X(A, B, T_1)

grant - S(A, T_2)

grant - S(B, T_2)

grant - X(B, T_1)

To finish the execution of both the transaction T_1 and T_2 :

- (1) T_1 is holding the lock on the data item A in an incompatible mode and this is required by T_2 .
- (2) T_2 --- and this is

None of T_1 & T_2 can finish their execution.
This situation is called ~~dead~~ DEADLOCK.

If the system is DEADLOCKED then it must rollback one of the transactions.

Ques. Which transaction we should choose??

Goal is to minimize the amount of rework.

Deadlock

Vs. Inconsistent state X

Incorrect result will be stored
in database.

Ques. Can we develop a locking protocol (policy)
such that the possibility of entering into
a deadlocked state is very very less ??